Evaluation findings

Sideline | Implementation snapshot: October 11, 2026

Question: Does a model using opportunity, efficiency and opponent context improve PPR forecasting over a recent average?

The model uses 24 lagged features and trains on 5,168 receiver appearances from 2023-2024. The eight-game baseline uses prior active/recorded appearances, including zero-event games where available.

2025 validation: 2,585 comparisons
Enhanced MAE: 4.218 PPR. Baseline MAE: 4.207 PPR.
Enhanced RMSE: 5.764. Baseline RMSE: 5.921.

2026 Weeks 1-4 diagnostic replay: 560 comparisons
Enhanced MAE: 4.504 PPR. Baseline MAE: 4.539 PPR.
Enhanced RMSE: 6.351. Baseline RMSE: 6.514.

Interpretation: The enhanced model does not improve 2025 average absolute error. It lowers RMSE in both evaluations, suggesting fewer or smaller large misses. The 2026 improvement is small. These results do not establish consistent superiority.

2026 outcomes had already been inspected, so this is a diagnostic replay rather than an untouched test. Current snapshots may contain later corrections. Forecasts assume the player appears and game/team context is known.

My roster: verified Week 1

Sideline | Implementation snapshot: October 11, 2026

Matthew Golden: projected 3.87, actual 15.5 PPR

Davante Adams: projected 13.19, actual 5.6 PPR

Emeka Egbuka: projected 6.28, actual 11.3 PPR

Luther Burden III: projected 9.57, actual 9.5 PPR

Marvin Harrison Jr.: projected 7.27, actual 4.3 PPR

Enhanced MAE: **5.456 PPR**. Original recorded-game-average baseline MAE: **7.424 PPR**. These are five user-confirmed league scores, which match nflverse. This small sample does not establish general accuracy.

The app favors Adams over Golden in this example, but Golden scores 15.5 versus Adams at 5.6. A useful comparison interface still needs to acknowledge forecast uncertainty.

Data checks and practical limits

Sideline | Implementation snapshot: October 11, 2026

The source audit covers 11,616 receiver appearances. Snap share is available for about 96% and drop statistics for about 77%. Missing model inputs use training medians. The interface labels unavailable values.

The automated suite contains 14 passing checks, including future-outcome leakage and invalid comparison inputs. This does not replace testing on another computer or testing with users.

Target share and red-zone share use explicit team denominators. Defense context uses previous passing EPA and WR PPR allowed. Recent trends enter as features instead of an arbitrary hot-streak bonus.

Routes, reliable pregame injuries and archived weather forecasts remain future work. Historical game-day weather would introduce information unavailable at prediction time.

Source: https://github.com/nflverse. Reproducible results: results/enhanced/evaluation.json and verified_week1.json. Source manifests and implementation details accompany the code.